GOOGLE MEET · AUTO-GENERATED TRANSCRIPT

OpenShift Cluster Kickoff — Hartwell Logistics × Red Hat

Monday, May 4, 2026 · 11:00 AM – 11:30 AM ET · Google Meet (meet.google.com/zpw-dnxh-kkr) · Recording & transcription enabled

Attendees (9)

Name	Role	Organization
Priya Venkatesh	Engagement Manager	Red Hat
Marcus Brennan	Senior Solutions Architect	Red Hat
David Okonkwo	Senior Consultant	Red Hat
Elena Voss	CTO	Hartwell Logistics
Raj Patel	Principal Architect	Hartwell Logistics
Liam O'Connor	Senior Engineer	Hartwell Logistics
Sofia Ramirez	Senior Engineer	Hartwell Logistics
Yuki Tanaka	Platform Engineer	Hartwell Logistics
Ben Chen	Backend Engineer	Hartwell Logistics

Transcript automatically generated by Google Meet. Speaker attribution and timestamps are based on voice recognition and may contain minor inaccuracies. Lightly cleaned for readability.

Transcript

Priya Venkatesh · 11:00:08

Okay — I think we have everybody. Good morning. Let me share the agenda while folks settle in.

Priya Venkatesh · 11:00:21

We have thirty minutes, nine of us, so I'm going to keep things moving. Quick intros, then scope and timeline from me, architecture and install from Marcus, day-2 and developer flow from David, training, and we'll save the last couple of minutes for open questions. Sound okay, Elena?

Elena Voss · 11:00:48

Sounds good. Thank you for the agenda — it's the right shape.

Priya Venkatesh · 11:00:56

Great. Let's do thirty-second intros, Red Hat side first. I'm Priya, engagement manager — I'm your single point of contact for anything that's not a technical question.

Marcus Brennan · 11:01:11

Marcus Brennan. I've been a solutions architect at Red Hat for about nine years, mostly on bare-metal OpenShift installs. I'll own the cluster design and the install.

David Okonkwo · 11:01:28

David Okonkwo, senior consultant. I'll be on the ground with your team for the day-2 ops piece — runbooks, monitoring, the developer workflow stuff.

Elena Voss · 11:01:42

Elena Voss, CTO. I'll be in and out of the working sessions but I wanted to be in the room for kickoff. Raj is the day-to-day owner.

Raj Patel · 11:01:55

Raj Patel, principal architect. I've been the one trading emails with Priya, and I'll be the technical decision-maker on our side.

Liam O'Connor · 11:02:08

Liam O'Connor, senior engineer. I lead our payments service team. Probably the loudest voice on the developer side.

Sofia Ramirez · 11:02:18

Sofia Ramirez, also senior engineer. I run the platform we have today, such as it is.

Yuki Tanaka · 11:02:28

Yuki Tanaka, platform engineering. Mostly here to ask uncomfortable questions about Operators.

Ben Chen · 11:02:38

Ben Chen, backend. I'll be one of the people actually pushing code to the cluster. I read the pre-read.

Priya Venkatesh · 11:02:48

Thank you, all. Quick scope and timeline before we hand to Marcus.

Priya Venkatesh · 11:02:58

What we are doing: standing up an OpenShift Container Platform Plus cluster on your bare metal in Columbus, with Advanced Cluster Management and Quay. The install window is two weeks starting May 18. Stabilization is four weeks after that. Day-2 handoff is at day forty-five — we renegotiated that from thirty in the SoW, Raj, you'll remember.

Raj Patel · 11:03:28

I do remember. Thanks for that.

Priya Venkatesh · 11:03:32

Production cutover target is end of Q2. That's the date everything flexes around. Marcus — over to you on the architecture.

Marcus Brennan · 11:03:44

Thanks, Priya. I'm going to skip the slides since you've all read the architecture overview, and just hit the three things people usually want to ask about.

Marcus Brennan · 11:04:00

First, the cluster shape. Three control plane nodes, twelve workers to start, room to add four more without re-cabling. Each worker is the same SKU as your current cluster, which simplifies your operations story and gives Sofia one less thing to keep in her head.

Sofia Ramirez · 11:04:22

Appreciated.

Marcus Brennan · 11:04:25

Second, install time. The install itself, agent-based, runs in about four hours once the prerequisites are in place. The reason I'm asking for two weeks is the prerequisites — DNS, load balancer rules, the VLAN cutover, the firewall reviews. That's the ninety percent of the calendar that isn't OpenShift. If your network team is ready when we land, we'll be done in week one and use week two for hardening.

Yuki Tanaka · 11:05:02

When you say 'agent-based,' is that the same agent installer that replaced UPI on baremetal, or is this something newer?

Marcus Brennan · 11:05:14

Same one. It's been the recommended path for bare metal for a while now. Generates the ISO, you boot the nodes, the cluster forms itself. We'll do it together so your team sees it end to end.

Yuki Tanaka · 11:05:32

Okay. I had read the older docs. Good to know.

Marcus Brennan · 11:05:38

Third — and this is the one Raj asked me to make sure I covered — how your developers get from 'I have a service' to 'it's running on the cluster.' On OpenShift you have three reasonable paths, and you do not have to pick just one.

Marcus Brennan · 11:06:02

Path A is your existing CI builds the image, pushes to Quay, ArgoCD syncs the manifest. That's what Sofia and Liam will recognize from today's setup. We'll keep this for the services that already work this way.

Marcus Brennan · 11:06:25

Path B is Source-to-Image, where the cluster builds the container from your source repo. It's the fastest way to get a developer from git push to running pod with no Dockerfile. It's also the path that produces the most consistent images across teams.

Marcus Brennan · 11:06:48

Path C is BuildConfigs with an explicit Dockerfile, for the cases where you need full control. We'll set up examples of all three in the first stabilization week so Ben and the team can pick what fits each service.

Ben Chen · 11:07:15

The thing I'm trying to figure out is — for our payments service, where the build is custom and we run a security scan as part of CI, is Path A the right fit? I don't think we can move that scan into the cluster easily.

Marcus Brennan · 11:07:36

Yes. Path A is right for payments. The rule of thumb is: if your CI is doing real work — security scans, SBOM generation, license checks — keep it in CI. S2I shines for the new green-field stuff where the build is just 'compile and package.'

Liam O'Connor · 11:07:58

Quick one — when we push to Quay, are we pushing to a Quay that lives in the cluster, or a separate Quay?

Marcus Brennan · 11:08:10

In the cluster. Quay is part of the Platform Plus install. We'll stand it up on the same nodes as the rest of the workload, with its own storage class for the registry. You get image scanning out of the box.

Liam O'Connor · 11:08:30

Good. That replaces a thing we've been meaning to fix for a year.

Priya Venkatesh · 11:08:36

Marcus, I'm going to gently move us along — David has the day-2 piece and I want to leave time for questions.

Marcus Brennan · 11:08:48

Yeah, handing off.

David Okonkwo · 11:08:52

Thanks. Day-2 ops, in plain English, is everything between 'install complete' and 'it's been running quietly for a year.' Patching, monitoring, on-call, certificate rotation, capacity planning, the boring stuff that decides whether the platform is loved or hated.

David Okonkwo · 11:09:18

What we're proposing is a forty-five-day window where Red Hat owns the operations practice — meaning Marcus and I are the ones in the monitoring tools, writing the runbooks, walking your platform owner through the failure modes. We're not on a pager rotation; that's what Premium Support is for. We're building the muscle memory.

Elena Voss · 11:09:50

On the Premium Support point — I want to confirm what I read in your email last week. If we have a Sev-1 at 2 AM on day twenty, the call goes to Premium Support, not to you?

David Okonkwo · 11:10:08

Correct. Premium Support is twenty-four-by-seven from day one. They're set up for that — we are not. What you'll find, though, is that by day twenty most of the calls won't be to Premium Support. They'll be to your own platform owner, because the runbook will tell them what to do.

Elena Voss · 11:10:32

That's the answer I wanted. Thank you.

David Okonkwo · 11:10:38

On developer self-service: this is the part that will feel different from your current setup. Today, when a team needs a namespace, they file a ticket with Sofia. We're going to replace that with a self-service flow — a developer fills out a short form, and a project request goes to a controller that creates the namespace, sets the quotas, applies the network policy, and notifies the team.

Sofia Ramirez · 11:11:12

How opinionated is the controller? I have ten teams and they all want different quotas.

David Okonkwo · 11:11:22

It's templated. You'll define maybe four or five tiers — small, medium, large, payments-special — and the form lets the team pick the tier with a short justification. The tier sets the quota. You can override per namespace if you need to, but the goal is that ninety-five percent of the time nobody overrides anything.

Sofia Ramirez · 11:11:48

Okay. That's manageable.

Yuki Tanaka · 11:11:54

Can I ask the Operators question now or is there a better moment?

Priya Venkatesh · 11:12:02

Now is fine.

Yuki Tanaka · 11:12:06

Some of our internal tools are deployed today as Helm charts. We've had a long internal debate about whether to write Operators for them or keep them as Helm. What do you tell people?

David Okonkwo · 11:12:24

Honest answer is, I don't write an Operator unless the thing has a real lifecycle — meaning it has state that has to be managed, upgrades that are non-trivial, or a reconciliation loop that matters. Helm is fine for stateless services. The trap is teams writing Operators because Operators feel modern, when a Helm chart and a GitOps sync would do the same job with a tenth of the maintenance.

Yuki Tanaka · 11:13:00

That is the answer I was hoping you would give.

David Okonkwo · 11:13:06

We can go deeper on it in the working session Thursday — happy to walk through three of your candidate services together and decide case by case.

Yuki Tanaka · 11:13:20

Yes please.

Priya Venkatesh · 11:13:24

David, can you cover the day-2 handoff before I take training?

David Okonkwo · 11:13:32

Sure. At day forty-five, your platform owner — and I'm going to push Raj to name that person before week one ends — runs a one-hour tabletop with us. We simulate three failure modes: a node fails, an image push fails, a certificate expires. If they can talk us through the runbook for all three, the handoff is done. If not, we extend by two weeks. We'd rather extend than declare victory early.

Raj Patel · 11:14:08

Sofia is the platform owner. We talked about it last week.

Sofia Ramirez · 11:14:14

I am the platform owner. I'd like that on the record.

Priya Venkatesh · 11:14:20

Noted, in the recording and the minutes. Training, then questions.

Priya Venkatesh · 11:14:30

Three pieces. First, three back-to-back DO180 cohorts in late May, two days each, instructor-led. That covers your developers — containerization fundamentals, the OpenShift CLI, deploying their first applications. Second, one DO288 cohort in mid-June for the senior folks — that's the Operator and platform path, three days. Liam and Sofia, that's you, and we have headroom for two more from your team if you want.

Liam O'Connor · 11:15:08

Yuki should be in DO288.

Yuki Tanaka · 11:15:14

Yes.

Priya Venkatesh · 11:15:18

Done. Third piece is two half-day workshops we'll tailor to your specific CI pipeline — taking the patterns Marcus described and applying them to two real services from your codebase. We'll pick the services with you in week two of the install.

Priya Venkatesh · 11:15:48

Cert vouchers are included for anyone who wants to take EX180 or EX288. Six EX180 vouchers and four EX288 are in scope; Raj, you've already committed six and two.

Raj Patel · 11:16:08

Yes. Voucher names went to your training team Friday.

Priya Venkatesh · 11:16:14

Perfect. We are at sixteen minutes — fourteen left. Open questions. Don't be polite.

Ben Chen · 11:16:28

I have one. The pre-read mentioned ArgoCD as part of the pattern. Today we have a Jenkins pipeline that does the deploy step. Are we expected to throw that out, or is there a graceful migration?

David Okonkwo · 11:16:48

Graceful. Day one, your Jenkins pipeline still runs and still deploys; we just point its kubectl apply at the new cluster. Over the first month, we move services one at a time onto the ArgoCD model — you'll feel the difference, because suddenly the deploy step in Jenkins becomes a git commit instead of a kubectl call. The cluster pulls instead of being pushed to. But we don't force the migration on day one.

Ben Chen · 11:17:30

Good. That was the answer that was going to determine whether I had to rewrite the payments pipeline this quarter.

Sofia Ramirez · 11:17:38

Mine is about backups. The current cluster has a backup story that I am not proud of. What's the right pattern on the new platform?

Marcus Brennan · 11:17:54

OADP — OpenShift API for Data Protection. It's the supported path. Backs up etcd, namespaces, persistent volumes to your object store of choice. We'll set it up against your existing S3 — your team already has buckets, no new procurement. Takes about half a day in the stabilization week.

Sofia Ramirez · 11:18:22

That is a much better answer than what I have today.

Liam O'Connor · 11:18:30

Question on the educational side. Beyond the certifications, are there resources you'd recommend for the engineers who finish DO288 and want to keep going? Especially around designing for the platform rather than just deploying to it.

Priya Venkatesh · 11:18:54

Two things. First, the Red Hat Developer site has a learning-paths section that's free, and the Application Development paths are actually good — not just marketing. Second, once we're in delivery I'd like to do a monthly thirty-minute office hour with David where your team brings real problems and we work through them. It's not in the SoW, but I've been doing it with another customer and it's the highest-leverage hour of my month.

Liam O'Connor · 11:19:30

Yes. Put me down for that.

Elena Voss · 11:19:34

Add me to the office hour distribution as well, even if I only show up half the time.

Priya Venkatesh · 11:19:42

Done.

Yuki Tanaka · 11:19:46

One more, on observability. We use Grafana and Prometheus today. What does OpenShift bring in the box, and where do we keep our stack versus replace it?

Marcus Brennan · 11:20:06

OpenShift ships with a Prometheus stack for cluster monitoring — node health, etcd, kubelet, the platform itself. You should not replace that. For your application metrics, you have two choices: use the user workload monitoring feature, which is the same Prometheus extended to your namespaces, or keep your own Grafana and Prometheus and federate. Most customers your size do the latter, because they already have dashboards they like.

Yuki Tanaka · 11:20:42

We have dashboards we like.

Marcus Brennan · 11:20:48

Then federate. We'll show you the pattern in the stabilization week.

Raj Patel · 11:20:54

I have a process question. We talked about working sessions later this week — what's actually on the calendar?

Priya Venkatesh · 11:21:04

Three sessions on the books, all on your calendar. Wednesday afternoon is the network and DNS deep-dive — Marcus and your network team. Thursday morning is the Operators and developer workflow session David mentioned. Friday morning is a backlog and sprint-zero session with me, Raj, and Sofia, where we'll carve up the next two weeks into something we can actually track.

Raj Patel · 11:21:38

Perfect.

Elena Voss · 11:21:44

I want to ask one strategic question before we close. What does success look like in twelve months — from your point of view, Priya, not from mine?

Priya Venkatesh · 11:22:00

Honest answer? In twelve months I want to be irrelevant. I want Sofia running the platform without thinking about us, your developers self-serving namespaces and not knowing it's OpenShift underneath, and the second cluster in Berlin standing up in two weeks instead of two months because your team did it and not us.

Priya Venkatesh · 11:22:32

If we get there and you only call us when you want to add a capability, the engagement worked. If you're calling us every week to keep the lights on, we built the wrong thing together.

Elena Voss · 11:22:52

That's the right answer. Thank you.

Raj Patel · 11:23:00

From our side — anything you need from us in the next two weeks to be ready for May 18?

Marcus Brennan · 11:23:12

Three things. The DNS records and the load balancer VIPs we sent in the prerequisites doc — we need those ready by May 11. The firewall change requests submitted by May 13. And someone from your network team in the Wednesday session this week.

Raj Patel · 11:23:36

All three are in flight. I'll confirm by Friday.

David Okonkwo · 11:23:42

And from me — Sofia, I'd like a thirty-minute walk-through of the current cluster's monitoring setup before Thursday so I'm not asking you basic questions in front of a room.

Sofia Ramirez · 11:23:58

I'll send a calendar hold.

Priya Venkatesh · 11:24:06

Anything else? I have about five minutes left on the clock.

Liam O'Connor · 11:24:14

Last one. Are we doing this in a way that I can show my engineers during the install? I'd like the team to see Marcus actually doing the install rather than getting the result handed to them.

Marcus Brennan · 11:24:32

Yes — that's how I prefer to do it. The install is in your datacenter, but the actual cluster bring-up is over a screen share that anyone on your team can join. I'll narrate. It is also somewhat boring after the first hour, fair warning.

Liam O'Connor · 11:24:54

We'll take the boring.

Elena Voss · 11:25:00

Good. I think we're in good shape. Priya, thank you for keeping us on time.

Priya Venkatesh · 11:25:10

Thank you, all. Recapping the actions in the chat: Raj confirms DNS, LB, and firewall by Friday; Sofia and David schedule the monitoring walk-through; Yuki, Liam, Sofia, and one more confirmed for DO288; office hour added to the engagement at no extra cost. Anything I missed?

Raj Patel · 11:25:38

Sofia is the named platform owner.

Priya Venkatesh · 11:25:42

Sofia is the named platform owner. Adding.

David Okonkwo · 11:25:48

I'll send the runbook template by end of day so the team can see what we're going to be filling in together.

Priya Venkatesh · 11:25:58

Perfect. See everyone Wednesday. Thank you.

Elena Voss · 11:26:06

Thank you.

Raj Patel · 11:26:10

Thanks, all.

Liam O'Connor · 11:26:14

Bye.

Sofia Ramirez · 11:26:16

Bye.

Yuki Tanaka · 11:26:18

Thanks.

Ben Chen · 11:26:22

Bye.

Marcus Brennan · 11:26:28

Priya, hold for one minute on my side?

Priya Venkatesh · 11:26:32
Yep. (Pause.) Okay, what's up?

Marcus Brennan · 11:26:40
Just want to flag — Yuki is going to be the most useful person on their team for us. He's already thinking about Operators the right way. Worth pairing him with David in the Thursday session more than the others.

Priya Venkatesh · 11:27:02
Noted. I'll add it to the prep doc.

David Okonkwo · 11:27:08
Agreed on Yuki.

Priya Venkatesh · 11:27:14
Anything else? Otherwise I'll let you go.

Marcus Brennan · 11:27:20
That's it.

David Okonkwo · 11:27:24
Talk Wednesday.

Priya Venkatesh · 11:27:28
Wednesday.

(End of recording — 11:27:42 ET)