

Hardware Pivot — Hartwell Logistics × Red Hat

Thursday, May 21, 2026 · 2:00 PM – 2:45 PM ET · Google Meet (meet.google.com/zpw-dnxh-kkr) ·
Recording & transcription enabled

Subject: Emergency replan — bare-metal hardware out of spec; decide pivot path and new prod cutover date.

Attendees (9)

Name	Role	Organization
Priya Venkatesh	Engagement Manager	Red Hat
Marcus Brennan	Senior Solutions Architect (on-site, Columbus DC)	Red Hat
David Okonkwo	Senior Consultant	Red Hat
Elena Voss	CTO	Hartwell Logistics
Raj Patel	Principal Architect	Hartwell Logistics
Liam O'Connor	Senior Engineer	Hartwell Logistics
Sofia Ramirez	Senior Engineer / Platform Owner	Hartwell Logistics
Yuki Tanaka	Platform Engineer	Hartwell Logistics
Ben Chen	Backend Engineer	Hartwell Logistics

Transcript automatically generated by Google Meet. Speaker attribution and timestamps are based on voice recognition and may contain minor inaccuracies. Lightly cleaned for readability.

Transcript

Priya Venkatesh · 14:00:11

Okay, I think we have everyone. Thanks for jumping on with two hours' notice. Marcus is dialing in from the data center floor in Columbus, so if his audio gets noisy, that's why.

Marcus Brennan · 14:00:32

Confirming. I'm in the row next to the cluster. I'll move if it gets loud.

Priya Venkatesh · 14:00:42

Forty-five minutes. I want to use the time to do four things: Marcus walks us through what we found, Sofia and Yuki tell us the network and storage piece from your side, we look at three pivot options together, and we land on a decision and a new schedule before we drop. Sound right, Elena?

Elena Voss · 14:01:14

That's the right shape. I want to leave this call with a number — what's the new prod cutover date — even if it's an estimate.

Priya Venkatesh · 14:01:28

Understood. You'll have one. Marcus, set the scene.

Marcus Brennan · 14:01:36

Short version: the install is paused. Long version: we found four things during pre-flight, and any one of them would have stopped us. I want to walk through them in order so we're all looking at the same problem.

Marcus Brennan · 14:01:58

Issue one. Six of the twelve worker nodes shipped with single-port 1-gigabit NICs instead of the dual-port 25-gigabit NICs in the bill of materials. We caught it Monday afternoon when I started cabling. The serial numbers match the order; the SKU on the cards does not.

Marcus Brennan · 14:02:30

Issue two. Four of those same six nodes are also on a baseboard management firmware version from late 2023 that won't accept the agent installer ISO over virtual media. We can flash them, but each flash is forty-five minutes and one node has bricked once already during the upgrade attempt.

Marcus Brennan · 14:02:58

Issue three. One control-plane node — the third one, the one in the middle — was DOA. Power supply isn't seating, the chassis throws a fault on boot. Our hands at the cage think it's a back-plane issue, not a power supply, which means it's a board RMA, not a parts swap.

Marcus Brennan · 14:03:24

Issue four, and this is the one I'm most worried about. The storage shelf for Ceph shipped with the previous-generation controller cards. The cards are functional — they speak the right protocol — but they cap out at about forty percent of the IOPS we sized for. We will run the cluster, we will not be happy with the cluster.

Elena Voss · 14:04:02

How is this possible? We did a hardware review with the vendor in March.

Marcus Brennan · 14:04:10

I have a theory. The BOM matches what you ordered. The picking and shipping doesn't. My guess is the vendor pulled from the wrong bin for the NICs and the controller cards, and the DOA node and the firmware are unrelated bad luck. I called my contact at the OEM yesterday and we have a ticket open. They are not arguing the facts.

Raj Patel · 14:04:42

Are they replacing or are they going to make us fight for it?

Marcus Brennan · 14:04:48

Replacing. The argument is going to be about how fast. They quoted me three to four weeks for the NIC swap kits and the new controller cards. The DOA chassis is a separate RMA and they're saying ten to fourteen days.

Sofia Ramirez · 14:05:12

I want to add one thing on the network side that the team noticed Tuesday. The two top-of-rack switches we put in for the new cluster are configured for the 25-gig profile already. If we drop 1-gig uplinks into them on the bad nodes, the switch will negotiate down but it'll throw a warning every reboot. I don't love that as a long-term picture, even if we ran it temporarily.

Yuki Tanaka · 14:05:48

And on storage — we did a quick fio benchmark on the shelf yesterday. It came in at thirty-eight percent of the spec sheet. So Marcus's number is right. We pulled the controller part numbers and they are previous-gen, confirmed.

Elena Voss · 14:06:14

Okay. So we are not arguing about whether we have a problem. Good. What are the options?

Priya Venkatesh · 14:06:24

Three real options on the table. I'll lay them out, then I want Marcus and David to push on each one, and then we decide. Option A is the conservative path: pause the install entirely, wait for the replacement hardware, restart in mid to late June, accept the schedule slip on prod. Option B is the aggressive path: install on the six good worker nodes and the two good control-plane nodes with a temporary three-control-plane workaround, run with the underspec storage, and swap hardware in flight as it arrives. Option C is the parallel path: stand up Red Hat OpenShift on AWS — ROSA — this week, run the developer enablement track and any non-production work there immediately, and let the bare-metal install settle when the hardware is right. The cluster you actually want in production gets built once, not twice.

Elena Voss · 14:07:38

Walk me through option C. We did not budget for AWS.

Priya Venkatesh · 14:07:46

Two answers. First, ROSA can be billed against existing AWS spend if you have an AWS bill, which I believe Hartwell does. Second, Red Hat can offer a sixty-day proof-of-value entitlement on the OpenShift side at no charge — the AWS infrastructure cost is yours, but the OpenShift subscription for those sixty days is on us, so long as we plan to migrate workloads to bare metal afterward, which we do.

Elena Voss · 14:08:24

What's the AWS infra cost, roughly?

Marcus Brennan · 14:08:30

For a four-worker dev-sized ROSA cluster running for sixty days, you're looking at roughly twelve to fifteen thousand dollars, plus or minus. Closer to twelve if we right-size and turn it off outside business hours.

Elena Voss · 14:08:54

That's an order of magnitude under the cost of slipping prod. Keep going.

David Okonkwo · 14:09:02

I want to push on option B before we get attached to anything. The reason I don't love installing on the partial hardware is the operational practice we're trying to build. If we run on degraded nodes for three weeks and then start swapping hardware, we are teaching Sofia and the team that the cluster is in flux, and we're writing runbooks against a configuration that doesn't exist anymore by July. Day-2 ops needs

a stable target.

Sofia Ramirez · 14:09:36

Agreed. From the operator's seat, I would rather start over once with the right hardware than rebuild trust in the cluster after three weeks of band-aids.

Liam O'Connor · 14:09:52

But I want to push on option A from the developer side. If we wait six weeks before any of us touches OpenShift, we lose the momentum from training. The DO180 cohorts are scheduled for next week. People will go to class, learn the platform, and then have nothing to apply it to until late June. That's a real cost.

Yuki Tanaka · 14:10:18

Plus, the four-week stabilization window in the SoW assumed we were running real workloads against the cluster, not just kicking tires. We need developers actually building things or we don't get any of the value of that window.

Ben Chen · 14:10:38

And our internal teams are blocked on a couple of services that were timed to the new cluster. I have a stack of containerization work I was going to start in June. If there's nothing to deploy to, that work doesn't start, and that's three or four weeks of engineering on the floor.

Elena Voss · 14:11:04

I'm hearing option C is the consensus. Anybody actually want to argue for A or B?

Marcus Brennan · 14:11:14

I'll steel-man option A briefly. The argument for pausing is that ROSA, even though it's the same OpenShift, isn't the same cluster. Anything we build there has to be lifted-and-shifted to bare metal in July. That migration is real work, and it's work we wouldn't have to do if we just waited.

Marcus Brennan · 14:11:42

But the migration is mostly redeploying applications against a new cluster, which is the same exercise we'd do anyway when prod cuts over. The training, the runbook foundation, the developer self-service patterns — none of that is wasted. The wasted work is small.

David Okonkwo · 14:12:04

And on B — I'll steel-man the install-on-degraded-hardware argument. The benefit is we never stop. The cluster exists, developers can deploy, we don't have to set up ROSA. The cost is what I said: we are debugging hardware while teaching people the platform, and the two get tangled. I've seen this before. The team blames OpenShift for problems that are actually a 1-gig NIC saturating, and the platform's reputation never recovers.

Sofia Ramirez · 14:12:42

That last point is the one. I would rather defend OpenShift than defend OpenShift-on-broken-hardware. Once the platform gets a reputation for being slow, it doesn't go away.

Raj Patel · 14:13:00

I'm with C. But I want to make sure we're being honest about what happens to bare metal. We don't quietly forget about it because ROSA is comfortable, do we?

Priya Venkatesh · 14:13:16

Good question, and the answer is no. The way I want to structure this is that ROSA is explicitly temporary, with a target migration date written into the schedule. If we don't put a date on it, you are right, it'll drift. So part of what I'll commit to today is a specific cutover date for moving off ROSA onto bare metal.

Elena Voss · 14:13:48

What's the new schedule? Marcus, give me the picture.

Marcus Brennan · 14:13:56

Three tracks running in parallel. Track one is ROSA: I can have a four-worker dev cluster up by end of next week, May 29. That's the platform your developers train against and start building on.

Marcus Brennan · 14:14:20

Track two is bare metal. We accept the OEM RMA timeline. Replacement NIC kits and storage controllers arrive between June 12 and 19. The DOA control-plane node arrives sooner, around June 1. I'd plan to be back on site June 22 for a full re-install — and this time it's a re-install, not a recovery from a partial. We throw away what we tried this week and start clean. Cluster is complete by July 3.

Marcus Brennan · 14:14:58

Track three is the migration. Two weeks for stabilization on bare metal — same four weeks we'd planned, but now compressed because the team has already been on ROSA and knows the platform. Then a two-week migration window where we move workloads off ROSA onto bare metal a service at a time. Prod cutover lands the week of July 20.

Priya Venkatesh · 14:15:34

So end of Q2 is gone. New target for prod is July 20. I want to say that out loud because it's the number Elena asked for. The schedule shifts right by roughly four weeks. That is real, and I'm not going to soften it.

Elena Voss · 14:16:00

Thank you for saying it directly. Four weeks is recoverable. Six weeks would have been a board conversation. What's the confidence level on the July 20 number?

Priya Venkatesh · 14:16:18

Honest answer is around seventy percent. The two things that move it are the OEM hitting their stated dates — they're usually within a week — and nothing else surprising us during the re-install. Marcus would put it slightly higher, I'd put it at seventy.

Marcus Brennan · 14:16:42

I'd put it at seventy-five. Priya is being conservative, which is her job.

Elena Voss · 14:16:50

Take the conservative number. What are the asks of us?

Priya Venkatesh · 14:16:58

Three things. One: we need a decision today on going to ROSA. I'd like to walk out of this call with a yes, so the AWS provisioning starts tomorrow. Two: we need someone from your finance team to approve the AWS spend — Raj, can you nominate the right person? Three: we need the OEM RMA paperwork

co-signed. Marcus has the draft.

Raj Patel · 14:17:32

Finance is Priya — actually our finance, also named Priya, confusingly. Priya Shah. I'll route to her this afternoon. She moves fast.

Elena Voss · 14:17:46

On the ROSA decision: yes. Provision it. I'll back-channel Priya Shah on the AWS approval so it doesn't sit in a queue.

Priya Venkatesh · 14:18:00

Decision logged. Thank you.

Liam O'Connor · 14:18:06

Question on the migration. When we move off ROSA onto bare metal in July — are the manifests portable? Is anything we write against ROSA going to need rewriting?

David Okonkwo · 14:18:24

Almost everything is portable. The two places to be careful are storage classes and load balancers. ROSA gives you AWS EBS by default; bare metal will be Ceph. So if a manifest hard-codes the storage class name, you'll have to change it. Same for load balancers — ROSA uses AWS NLBs, bare metal uses MetalLB.

David Okonkwo · 14:18:58

What we'll do is set up the manifests with parameterized storage class and ingress from day one, so the change is a value file swap, not a rewrite. We can show the team the pattern in the first week of training.

Liam O'Connor · 14:19:20

That works.

Sofia Ramirez · 14:19:24

What about day-2 ops? The runbooks we write against ROSA are going to be different from the ones we need on bare metal. Networking is different, storage is different, the install is different.

David Okonkwo · 14:19:42

Genuinely good point, and I'd rather not pretend it isn't. The runbooks split into two categories. The application-level stuff — deploying, scaling, ingress, image management — is identical between ROSA and bare metal because OpenShift is OpenShift. The infrastructure-level stuff — node management, storage backend, network failure modes — is genuinely different.

David Okonkwo · 14:20:18

What I'd propose is we write the application runbooks now against ROSA, since they transfer cleanly. We write the infrastructure runbooks in late June against bare metal, where they actually apply. We don't try to do both at once.

Sofia Ramirez · 14:20:42

That's the right split. I was worried we'd be writing two versions of everything.

Yuki Tanaka · 14:20:52

Question on the Operators conversation we had at kickoff. Does ROSA change any of that? Same Operators, same lifecycle?

Marcus Brennan · 14:21:08

Same Operators, same OperatorHub, same lifecycle manager. The only place ROSA differs is that some Operators have AWS-specific variants — for example, the AWS Load Balancer Operator only applies on ROSA. We'd just not install those. Everything else, including any Operator you write yourselves, is identical.

Yuki Tanaka · 14:21:34

Good. Then I don't have to throw out anything we discussed.

Ben Chen · 14:21:42

What about Quay? The image registry. Where does that live during the ROSA period?

Marcus Brennan · 14:21:54

Two options. We can stand up Quay on ROSA, which works fine, or we can use the integrated OpenShift image registry, which is lighter weight and good enough for sixty days. I'd recommend the integrated registry on ROSA, then a proper Quay deployment on bare metal in late June. Less moving parts during the temporary window.

Ben Chen · 14:22:24

Fine. I just wanted to make sure I wasn't going to be pushing images into a void.

Priya Venkatesh · 14:22:32

Let me get the schedule changes captured in the chat so we have them in writing. I'll read out and people correct me if I'm wrong.

Priya Venkatesh · 14:22:48

Original: install May 18, stabilization through mid-June, prod cutover end of Q2. New: ROSA stand-up May 26 to 29. Developer training proceeds as scheduled — DO180 cohorts next week, DO288 mid-June. Bare-metal hardware replacements in by June 19. Bare-metal re-install starts June 22. Bare-metal cluster complete July 3. Stabilization and migration July 6 through 17. Prod cutover week of July 20. Day-2 handoff to Sofia at day forty-five from cutover, so early September.

Raj Patel · 14:23:38

That's right. One add — the working sessions we had on the calendar for next week, we keep those, but we point them at ROSA instead of bare metal, yes?

David Okonkwo · 14:23:54

Yes. Same agendas, different cluster. The Thursday Operators session especially still works exactly as planned.

Elena Voss · 14:24:08

I want to ask about the cost picture. Are we adding to the SoW?

Priya Venkatesh · 14:24:16

On the consulting side, no — the work is the same shape, just shifted in time. Marcus's on-site days count what they would have counted, we don't double-bill the install. On the platform side, the OpenShift

entitlement on ROSA is covered by the proof-of-value I mentioned, so no charge from us. The AWS infrastructure spend is yours and is the twelve to fifteen thousand range.

Priya Venkatesh · 14:24:54

What I'll do is send a one-page change order today that captures the new dates, the ROSA bridge, and the cost picture. No new money on the consulting side, just paperwork to keep procurement happy.

Raj Patel · 14:25:18

Procurement will appreciate the paperwork.

Elena Voss · 14:25:24

What do I need to communicate up — to my board, to my CEO?

Priya Venkatesh · 14:25:32

I'd frame it three ways. One: the engagement is on track on everything we control. Two: the hardware vendor missed spec, we have a real RMA in motion, and the OEM is being responsive. Three: rather than wait, we are using a temporary AWS bridge so the team starts building immediately, which protects the engineering momentum. Net effect on the date is four weeks. Net effect on the team's productivity is roughly zero, because they start on ROSA instead of waiting.

Elena Voss · 14:26:14

That's a story I can tell.

Priya Venkatesh · 14:26:20

If it helps I can write that up as a one-page memo for you by tomorrow morning.

Elena Voss · 14:26:30

Yes please. Send it to Raj and me both.

Marcus Brennan · 14:26:38

I want to flag one risk we haven't named. If ROSA is too comfortable and the team gets used to managed OpenShift, the transition to bare metal in July is going to feel like a step down — operationally, ROSA is easier. We should set expectations now that bare metal is more powerful but more work.

Sofia Ramirez · 14:27:08

I appreciate that. I'll start the team's mental model in the right place. ROSA is the rental car, bare metal is the car we own.

Liam O'Connor · 14:27:22

I like that framing.

Yuki Tanaka · 14:27:28

What about ACM and the multi-cluster story? That was one of the reasons we picked Platform Plus. Does it still apply during the ROSA period?

Marcus Brennan · 14:27:42

Yes, and it's actually a small bonus. We can stand up ACM on the ROSA cluster as the management hub, then when bare metal comes online in July, we register it as a managed cluster. By the time Berlin rolls around in 2027, you've already operated ACM across two clusters. The Berlin install becomes a managed-cluster registration instead of a from-scratch install.

Yuki Tanaka · 14:28:18

That's actually a benefit of this whole situation.

Marcus Brennan · 14:28:26

I wouldn't have gone looking for it, but yes.

Elena Voss · 14:28:34

Anything else I need to weigh in on before I drop?

Priya Venkatesh · 14:28:42

One thing. The DO180 cohort starts Tuesday. We want it to run as scheduled. The ROSA cluster won't be ready until end of next week, but the first day of DO180 is in our training environment anyway, not yours. So the cohort proceeds. By the time they get to the hands-on portion in week two, ROSA is up. No change there.

Elena Voss · 14:29:14

Good. I'm going to drop. Raj, you have my support on whatever you decide in the rest of the call. Thank you all — and Marcus, thank you for being on top of this. This could have been a much messier conversation.

Marcus Brennan · 14:29:32

Thank you, Elena. We'll get there.

Elena Voss · 14:29:38

Bye.

Raj Patel · 14:29:40

Okay, we have about fifteen minutes. Let me use it to lock down the working items on our side.

Raj Patel · 14:29:50

Sofia, you own the day-2 runbook split — application-level on ROSA, infrastructure-level on bare metal in June. Yuki, you own the ACM hub setup with Marcus, since that becomes the multi-cluster spine. Liam and Ben, you two are my eyes on the developer experience — I want to know within the first week on ROSA whether the self-service namespace flow actually works for your teams or whether it needs adjusting before we lift it to bare metal.

Liam O'Connor · 14:30:24

Got it. We'll start using it as soon as it's up.

Ben Chen · 14:30:32

Same. I have two services I want to be the guinea pigs.

Sofia Ramirez · 14:30:40

Confirming I own the runbook split. I'd like to do the first drafts inline with David rather than handing him finished docs to review.

David Okonkwo · 14:30:54

Inline is better. I'll set up a recurring Tuesday-Thursday thirty-minute window through June for runbook drafting.

Yuki Tanaka · 14:31:08

On ACM — Marcus, when do you want to start?

Marcus Brennan · 14:31:14

Once ROSA is live, week of June 1. ACM install is a half-day. We'll do it together.

Yuki Tanaka · 14:31:24

Calendar hold incoming.

Raj Patel · 14:31:30

Priya, on the change order — when can we expect it?

Priya Venkatesh · 14:31:38

End of day tomorrow. I want to circulate it internally first so I'm not sending you something my own delivery practice hasn't blessed.

Raj Patel · 14:31:54

Reasonable.

Priya Venkatesh · 14:31:58

The other thing I owe you is the memo for Elena. That goes out tomorrow morning, before the change order.

Raj Patel · 14:32:10

Both of those are useful. What about the schedule view — do you have a Gantt or are we going off the chat readout?

Priya Venkatesh · 14:32:24

I'll send a one-page schedule view alongside the change order. Same dates I read out, just in picture form. Easier for Elena to share.

Marcus Brennan · 14:32:38

I'd add one more deliverable from my side — a short post-mortem on the hardware issue itself, what we found, what the OEM is doing, what we're going to do differently on the receive process next time. That's for your team's records.

Sofia Ramirez · 14:33:00

Yes please. I want that for our procurement team — they need to tighten how we accept hardware deliveries.

Raj Patel · 14:33:14

Anything else? David?

David Okonkwo · 14:33:20

One. I want to make sure the developer training on ROSA includes the migration pattern. We don't wait until July to teach 'this is how you parameterize a manifest so it moves between clusters.' We teach it in week one of DO180, so the team writes the right shape from the start.

Liam O'Connor · 14:33:48

That should be the default anyway. Even outside this situation, manifests should be portable.

David Okonkwo · 14:33:58

Correct. We're just going to be more explicit about it.

Yuki Tanaka · 14:34:08

On the Operators session Thursday — do we still do it the way we planned, or do we adjust the candidate services because they're being deployed to ROSA first?

David Okonkwo · 14:34:24

Same candidates, same conversation. The Helm-versus-Operator decision is about the workload, not about the cluster underneath. If we decide a service should be an Operator on bare metal, it should be an Operator on ROSA too.

Yuki Tanaka · 14:34:46

Good. I have my list ready.

Raj Patel · 14:34:54

Last thing from me — the Berlin team. They were going to pick up the second cluster work in early 2027 based on what we did here. Does any of this change what we tell them?

Priya Venkatesh · 14:35:14

It actually helps them. By the time Berlin starts, you'll have two clusters in production — bare metal Columbus and the ROSA experience. They can choose which model they want, and they can see the cost and operational tradeoffs from your real numbers. Better dataset than they would have had.

Raj Patel · 14:35:42

Useful. I'll tell them to stay tuned.

Marcus Brennan · 14:35:50

I'll be on site through Friday wrapping up the cabling that we did get right and documenting what's connected to what. Even with the hardware swap coming, the rack work isn't wasted — the switches stay, the cabling diagram stays, the power layout stays.

Sofia Ramirez · 14:36:14

Good. I'd like to see the cabling diagram when it's done — we never had a real one for the old cluster and that's bitten us.

Marcus Brennan · 14:36:26

You'll have it Friday.

Priya Venkatesh · 14:36:32

Let me read out the action items so I can drop them into the follow-up email tonight.

Priya Venkatesh · 14:36:42

Me: change order and schedule view by end of day tomorrow; memo for Elena tomorrow morning; ROSA provisioning starts tomorrow. Marcus: post-mortem on the hardware issue, on-site through Friday, cabling diagram Friday, RMA paperwork co-signature this week. David: runbook split confirmed, recurring Tues/Thurs drafting sessions through June, Operators session Thursday proceeds as planned. Raj: route AWS approval to Priya Shah this afternoon. Sofia: own day-2 runbook split, partner with David. Yuki: own ACM hub setup with Marcus starting June 1. Liam and Ben: live test of the self-service namespace flow on ROSA in the first week of June, report back.

Raj Patel · 14:37:38

All accurate.

Priya Venkatesh · 14:37:42

Good. We have about seven minutes left. I want to use them to make sure nobody's leaving the room with an unanswered question. What's bothering anyone?

Sofia Ramirez · 14:38:00

Honest answer — I'm worried about the optics internally. People on the engineering floor are going to hear 'install is paused' and form an opinion. How do we communicate this so it doesn't land as 'this project is in trouble'?

Priya Venkatesh · 14:38:22

That's exactly the kind of thing the memo to Elena should help with. If she has the framing, it cascades down. I'll make sure the memo is also useful as a paragraph you can send to your team. Raj, you have license to share it.

Raj Patel · 14:38:46

I'll send it Friday with my own framing on top.

Liam O'Connor · 14:38:54

I'm okay. I think the parallel approach is the right call and it gives us something to do.

Yuki Tanaka · 14:39:02

I'm fine. I'm actually more excited than I was an hour ago — the ACM-across-two-clusters thing is a better learning experience than the original plan.

Ben Chen · 14:39:18

Same. I'd rather start on ROSA next week than wait six weeks for anything to be real.

David Okonkwo · 14:39:30

From my side, I'm comfortable. The thing I want everyone to take from today is that we hit a hardware problem and we redesigned the engagement in forty minutes. That's the muscle the platform needs to be able to flex on day-2 too. Same pattern.

Marcus Brennan · 14:39:58

Agreed. And — I'll say this directly — the reason we could do that is that the kickoff was the right shape. Everyone in this room knew what the other people did. That's the only reason this call took forty-five minutes instead of three.

Raj Patel · 14:40:24

I appreciate you saying it. Priya runs a good kickoff.

Priya Venkatesh · 14:40:32

Thank you. Okay — we have five minutes back. I'll let everyone have their afternoon. Action items hit your inbox tonight, change order and memo tomorrow, schedule view tomorrow afternoon. Same group reconvenes — Raj, can we put a thirty-minute check-in on the calendar for next Thursday May 28, just to make sure ROSA stand-up is on track?

Raj Patel · 14:40:58

Yes. I'll send the invite when we drop.

Priya Venkatesh · 14:41:08

Perfect. Thank you everyone.

Sofia Ramirez · 14:41:14

Thanks all.

Liam O'Connor · 14:41:16

Bye.

Yuki Tanaka · 14:41:18

Thanks.

Ben Chen · 14:41:20

Bye.

Raj Patel · 14:41:24

Marcus — quick before you drop. Do you need anything from us this week beyond what we already covered?

Marcus Brennan · 14:41:36

One small thing. I'd like Sofia's preferred format for the cabling diagram. Visio? Draw.io? Lucid? I can produce any of them but it's not worth doing twice.

Sofia Ramirez · 14:41:54

Lucid. We have a team license.

Marcus Brennan · 14:42:00

Lucid it is.

Raj Patel · 14:42:06

Then we're good. Marcus, get some dinner. You've been in that data center for four days.

Marcus Brennan · 14:42:14

Working on it.

Priya Venkatesh · 14:42:18

Marcus — I'll call you after this. Thanks all. Talk Thursday.

David Okonkwo · 14:42:26

Talk Thursday.

Raj Patel · 14:42:30

Bye.

(End of recording — 14:42:48 ET)